

**RATHINAM COLLEGE OF ARTS & SCIENCE (AUTONOMOUS)**  
**Coimbatore-641021**

**DEPARTMENT OF COMPUTER SCIENCE**

**Core Practical – MACHINE LEARNING Lab Manual**



**Lab Manual for the Academic Year 2023-24**

(in accordance with Computer Science syllabus)

**SUBJECT : Machine Learning Lab**

**STREAM : B.Sc. [CS/CT/IT/AIML]/BCA**

Staff Incharge

H.O.D

## **MACHINE LEARNING LAB MANUAL**

1. Statistical Analysis of Ozone Level and Weather Factors-File name-Lab1.R
2. Multiple Linear Regression and Diagnostic of Environment Data-File name-Lab2.R
3. Linear Regression analysis of Ozone Level-File name-Lab3.R
4. Linear Regression Analysis of Study, Sleep Hours and Exam scores-File name-linear.R
5. Linear Regression and Model Evaluation with Salary Data-File name- income.R
6. Alcohol Consumption Analysis using Scatter Plot- File name- al.R
7. Alcohol Consumption Analysis using Bar-Graph-File name- al2.R

# 1. Statistical Analysis of Ozone Levels and Weather Factors

Dataset for 1st 3 programs (Lab1, Lab2, and Lab3):

S. No.	Ozone	Solar R	Wind	Temp	Month	Day
1	41	190	7.4	67	5	1
2	36	118	8	72	5	2
3	12	149	12.6	74	5	3
4	18	313	11.5	62	5	4
5	27	192	14.3	56	5	5
6	28	193	14.9	66	5	6
7	23	299	8.6	65	5	7
8	19	99	13.8	59	5	8
9	8	19	20.1	61	5	9
10	24	194	8.6	69	5	10
11	7	152	6.9	74	5	11
12	16	256	9.7	69	5	12
13	11	290	9.2	66	5	13
14	14	274	10.9	68	5	14
15	18	65	13.2	58	5	15
16	14	334	11.5	64	5	16
17	34	307	12	66	5	17
18	6	78	18.4	57	5	18
19	30	322	11.5	68	5	19
20	11	44	9.7	62	5	20

## Aim:

To understand how weather conditions influence ozone levels through statistical analysis.

## Algorithm:

### Step 1: Install and Load Required Packages

- Install and load the necessary R packages for data manipulation, visualization, and statistical analysis.

### Step 2: Create the Data Frame

- Create a data frame called "Data\_Frame" containing data from the given table.

### Step 3: Summarize the Data

- Use the `summary` function to summarize the "Data\_Frame."

### Step 4: Linear Regression

- Fit a linear regression model, taking "Ozone" as the dependent variable and using multiple independent variables from the dataset.

#### Step 5: Predict Ozone Level for the 21st Day

- Create new data for the 21st day's factors and predict the ozone level using the fitted model.

#### Step 6: Autocorrelation Analysis

- Find the autocorrelation of the error produced from the fitted line.

#### Step 7: Multicollinearity Analysis

- Analyze multicollinearity among independent variables and identify suitable solutions to remove multicollinearity.

#### Step 8: Equal Variance Check

- Find the variance among error terms and comment on the equal variance among error terms in the output.

#### Step 9: Autocorrelation Test

- Estimate the presence of autocorrelation using the Durbin–Watson test statistic.

#### **Code:**

##### **# Install and load the required packages**

```
install.packages("dplyr")  
install.packages("ggplot2")  
install.packages("car")# Load required libraries  
library(dplyr)  
library(ggplot2)  
library(car)# Companion to applied regression
```

##### **# Create the data frame for the given table**

```
Data_Frame <- data.frame(  
  S.No. = 1:20,  
  Ozone = c(41, 36, 12, 18, 27, 28, 23, 19, 8, 24, 7, 16, 11, 14, 18, 14, 34, 6, 30, 11),  
  Solar_R = c(190, 118, 149, 313, 192, 193, 299, 99, 19, 194, 152, 256, 290, 274, 65, 334,  
307, 78, 322, 44),  
  Wind = c(7.4, 8, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9.7, 9.2, 10.9, 13.2, 11.5, 12,  
18.4, 11.5, 9.7),  
  Temp = c(67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64, 66, 57, 68, 62),
```

```
Month = c(5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5),  
Day = 1:20  
)
```

**# Task 1: Summarize the above table in R**

```
summary(Data_Frame)
```

**# Task 2: Find the linear regression line on the given table taking ozone as the dependent variable**

```
model <- lm(Ozone ~ Solar_R + Wind + Temp + Month + Day, data = Data_Frame)
```

```
summary(model)
```

**# Task 3: Predict the 21st day's ozone level in the air with given factors**

```
new_data <- data.frame(  
  Solar_R = 100,  
  Wind = 15,  
  Temp = 70,  
  Month = 5,  
  Day = 21  
)
```

```
predicted_ozone <- predict(model, newdata = new_data)
```

```
print(predicted_ozone)
```

**# Task 4: Find the autocorrelation of the error produced from the fitted line**

```
residuals <- residuals(model)
```

```
acf(residuals)
```

**# Task 5: Analyze multicollinearity among independent variables and find a suitable solution to remove multicollinearity**

```
alias_table <- alias(model)
```

```
print(alias_table)
```

**# Task 6: Find the variance among error terms and comment on the equal variance among error terms in the output**

```
plot(model, which = 1)
```

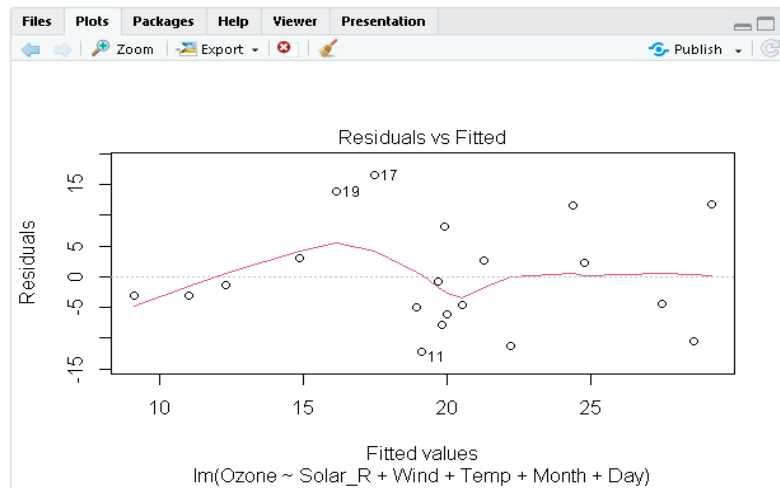
### # Task 7: Estimate the presence of autocorrelation using the Durbin-Watson test statistic

```
dw_test <- durbinWatsonTest(model)
```

```
dw_statistic <- dw_test$statistic
```

```
print(dw_statistic)
```

### Output:



## **2. Multiple Linear Regression and Diagnostic of Environment Data**

### **Aim:**

To assess the impact of multiple environmental factors on a specific environmental variable using multiple linear regression

### **Algorithm:**

#### **Step 1: Load Required Libraries**

In this step, we load the necessary libraries: ``stats`` for statistical functions, ``lmtest`` for additional linear regression testing, and ``car`` for regression diagnostics.

#### **Step 2: Create the Data Frame**

In this step, a data frame named ``Data_Frame`` is created, containing various environmental variables such as Ozone, Solar\_R, Wind, Temp, Month, and Day.

#### **Step 3: Estimate the Regression Model**

Multiple linear regression is performed with Ozone as the dependent variable and Solar\_R, Wind, and Temp as predictors. The ``lm`` function is used to create the regression model.

#### **Step 4: View Model Summary**

This step displays a summary of the regression model, including coefficients, standard errors, t-values, p-values, and various statistics for model evaluation.

#### **Step 5: Analyze Significance of Regression Coefficients**

The significance of regression coefficients is assessed using analysis of variance (ANOVA). The results are displayed.

#### **Step 6: Evaluate Model Fit**

R-squared and adjusted R-squared values are calculated to evaluate the goodness of fit of the regression model.

#### **Step 7: Diagnostic Checking**

Cook's Distance and the Press Statistic are computed for diagnostic checking of the regression model.

#### **Step 8: Post-Model Statistical Testing**

This step includes various post-model statistical tests to ensure a better model fit and accurate predictions, including tests for heteroscedasticity, autocorrelation, multicollinearity, and residual analysis.

#### **Step 9: Normality Testing on Error Terms**

A Shapiro-Wilk test is conducted to assess the normality of the error terms in the fitted model.

**Code:**

```
install.packages("stats")
```

```
install.packages("lmtest")
```

```
install.packages("car")# Load required libraries
```

```
library(stats)
```

```
library(lmtest)
```

```
library(car)
```

**# Create the dataframe**

```
Data_Frame <- data.frame(
```

```
  S_No = 1:20,
```

```
  Ozone = c(41, 36, 12, 18, 27, 28, 23, 19, 8, 24, 7, 16, 11, 14, 18, 14, 34, 6, 30, 11),
```

```
  Solar_R = c(190, 118, 149, 313, 192, 193, 299, 99, 19, 194, 152, 256, 290, 274, 65, 334, 307, 78, 322, 44),
```

```
  Wind = c(7.4, 8, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9.7, 9.2, 10.9, 13.2, 11.5, 12, 18.4, 11.5, 9.7),
```

```
  Temp = c(67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64, 66, 57, 68, 62),
```

```
  Month = c(5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5),
```

```
  Day = 1:20
```

```
)
```

**# 1. Estimate appropriate regression line with suitable predictors.**

**# We will perform multiple linear regression using Ozone, Solar\_R, Wind, and Temp as predictors.**

```
model <- lm(Ozone ~ Solar_R + Wind + Temp, data = Data_Frame)
```

```
summary(model) # View regression coefficients and statistics
```

**# 2. Estimate the significance of regression coefficients using ANOVA and compare with F and partial t test.**

```
anova_result <- anova(model)
```

```
anova_result
```



### **# 3. Model fit using R Square and Adjusted R square values.**

```
rsquared <- summary(model)$r.squared  
adj_rsquared <- summary(model)$adj.r.squared  
rsquared  
adj_rsquared
```

### **# 4. Estimate Cook Statistic and Press Statistic for diagnostic checking**

```
cooks_d <- cooks.distance(model)  
press_statistic <- sum((resid(model)/(1 - cooks_d))^2)  
cooks_d  
press_statistic
```

### **# 5. Post model statistical testing for better fit and error-free prediction.**

#### **# a. Breusch-Pagan Test for Heteroscedasticity**

```
bptest(model)
```

#### **# b. Durbin-Watson Test for Autocorrelation**

```
dwtest(model)
```

#### **# c. VIF (Variance Inflation Factor) for Multicollinearity**

```
vif(model)
```

#### **# d. Residual Analysis and Plots**

##### **# Check for linearity and homoscedasticity in residuals**

```
plot(model, which = 1)
```

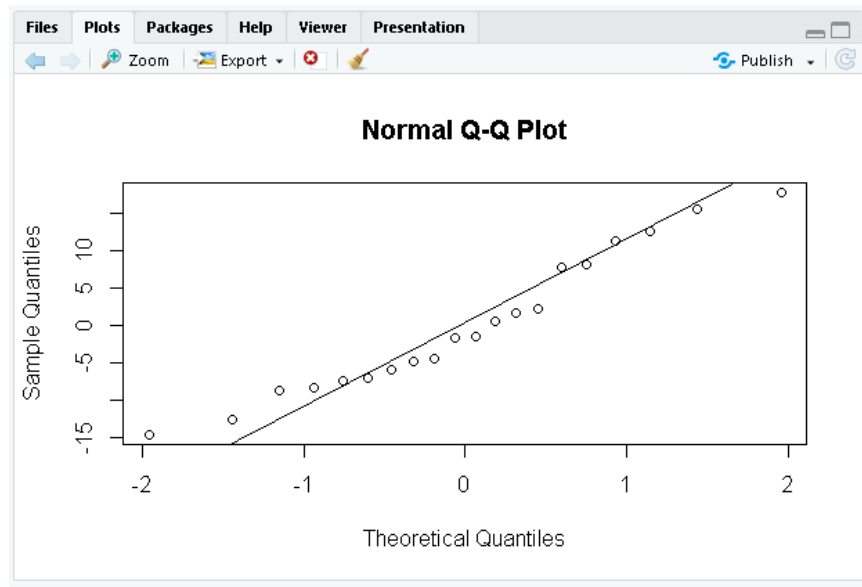
##### **# Plot a histogram of residuals and quantile-quantile (Q-Q) plot**

```
hist(resid(model))  
qqnorm(resid(model))  
qqline(resid(model))
```

### **# 6. Normality testing on error terms of the fitted model.**

```
shapiro.test(resid(model))
```

**Output:**



### 3. Linear Regression Analysis of Ozone Level

**Aim:**

To analyze and model ozone levels using linear regression

**Algorithm:**

**Step 1: Install and Load Required Packages**

- Install and load the necessary R packages for data visualization and variance inflation factor (VIF) calculation.

**Step 2: Create the Data Frame**

- Create a data frame named "Data\_Frame" containing information related to ozone levels and weather-related factors.

**Step 3: Plot Residuals Versus Fitted Values**

- Fit a linear regression model and plot residuals versus fitted values.

**Step 4: Plot Residuals Versus Observed Values**

- Plot residuals versus observed values.

**Step 5: Plot Observed Versus Fitted Values**

- Plot observed versus fitted values.

**Step 6: Find Maximum Leverage Value**

- Calculate and print the maximum leverage value and the corresponding observation index.

**Step 7: Interpret Residual Summary**

- Output the summary of the linear regression model to interpret the residual statistics.

**Step 8: Calculate VIF (Variance Inflation Factor)**

- Calculate the VIF values to assess multicollinearity among independent variables.

**Code:**

```
install.packages("ggplot2")
```

```
install.packages("car")
```

```
# Load required packages
```

```
library(ggplot2) # For plotting
```

```
library(car)    # For VIF calculation
```

### **# Create the dataframe**

```
Data_Frame <- data.frame(  
  S_No = 1:20,  
  Ozone = c(41, 36, 12, 18, 27, 28, 23, 19, 8, 24, 7, 16, 11, 14, 18, 14, 34, 6, 30, 11),  
  Solar_R = c(190, 118, 149, 313, 192, 193, 299, 99, 19, 194, 152, 256, 290, 274, 65, 334,  
307, 78, 322, 44),  
  Wind = c(7.4, 8, 12.6, 11.5, 14.3, 14.9, 8.6, 13.8, 20.1, 8.6, 6.9, 9.7, 9.2, 10.9, 13.2, 11.5, 12,  
18.4, 11.5, 9.7),  
  Temp = c(67, 72, 74, 62, 56, 66, 65, 59, 61, 69, 74, 69, 66, 68, 58, 64, 66, 57, 68, 62),  
  Month = c(5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5),  
  Day = 1:20  
)
```

### **# 1. Plot residual versus Fitted values using plot command**

```
model <- lm(Ozone ~ Solar_R + Wind + Temp, data = Data_Frame)  
residuals <- residuals(model)  
fitted_values <- fitted(model)  
plot(fitted_values, residuals, main = "Residuals vs. Fitted Values", xlab = "Fitted Values",  
ylab = "Residuals")
```

### **# 2. Plot residual versus Observed using Plot command**

```
plot(Data_Frame$Ozone, residuals, main = "Residuals vs. Observed", xlab = "Observed  
Values", ylab = "Residuals")
```

### **# 3. Plot observed versus and fitted values using plot command**

```
plot(fitted_values, Data_Frame$Ozone, main = "Observed vs. Fitted Values", xlab = "Fitted  
Values", ylab = "Observed Values")
```

### **# 4. Find out the leverage value in the fitted values using which.max command.**

```
leverage_values <- hatvalues(model)  
max_leverage_index <- which.max(leverage_values)  
max_leverage_value <- leverage_values[max_leverage_index]
```

```
cat("Maximum Leverage Value:", max_leverage_value, "\n")
```

```
cat("Observation Index with Maximum Leverage:", max_leverage_index, "\n")
```

**# 5. Interpret the residual summary from the lm( ) command.**


```
summary(model)
```

**# 6. Find out the VIF Variance Inflation Factor values using inbuilt function available in R.**

```
vif_values <- vif(model)
```

```
print(vif_values)
```

**Output:**



#### 4. Linear Regression Analysis of Study Hours, Sleep Hours, and Exam Scores

##### Dataset:

S.No.	study_hours	sleep_hours	examscore
1	2	7	70
2	3	6	80
3	5	5	90
4	6	6	85
5	4	8	75
6	2	5	65
7	4	5	85
8	8	6	65
9	6	8	55
10	3	4	65

##### Aim:

To explore and quantify the relationships between study hours, sleep hours, and exam scores through linear regression analysis.

##### Algorithm:

Step 1: Load Required Libraries

In this step, we load the `ggplot2` library for data visualization.

Step 2: Create a Data Frame

Here, a data frame named `Data\_Frame` is created. It contains three variables: `study\_hours`, `sleep\_hours`, and `examscore`.

Step 3: Print the Data Frame (Optional)

This step displays the contents of the `Data\_Frame` data frame.

Step 4: Explore the Data (Optional)

The `summary` function provides basic statistics about the data, such as mean, median, and quartiles. This step is optional but can help you understand the data better.

Step 5: Perform Linear Regression

In this step, a linear regression model is created. The dependent variable is `examscore`, and the independent variables are `study\_hours` and `sleep\_hours`. The model is stored in the `model` object.

Step 6: Summarize the Model

This step provides a summary of the linear regression model, including coefficients, standard errors, t-values, and p-values. It helps you understand the model's goodness of fit.

### Step 7: Predict Student Marks Using the Model

Here, a new data frame named `new\_data` is created with values for `study\_hours` and `sleep\_hours`. The `predict` function is used to make predictions based on the model, and the results are stored in `predicted\_marks`.

### Step 8: View the Predicted Marks

This step displays the predicted exam scores based on the input data.

### Step 9: Plot the Data and the Regression Line

#### **Code:**

#### **# Load required libraries**

```
library(ggplot2)
```

#### **# Create a data frame**

```
Data_Frame <- data.frame (  
  study_hours = c(2,3,5,6,4),  
  sleep_hours = c(7,6,5,6,8),  
  examscore = c(70,80,90,85,75)  
)
```

#### **# Print the data frame**

```
Data_Frame
```

#### **# Explore the data (optional)**

```
summary(Data_Frame)
```

#### **# Perform linear regression**

```
model <- lm(examscore ~ study_hours + sleep_hours, data = Data_Frame)
```

#### **# Summarize the model**

```
summary(model)
```

#### **# Predict student marks using the model**

```
new_data <- data.frame(study_hours = c(4), sleep_hours = c(7))
```

```
predicted_marks <- predict(model, newdata = new_data)
```

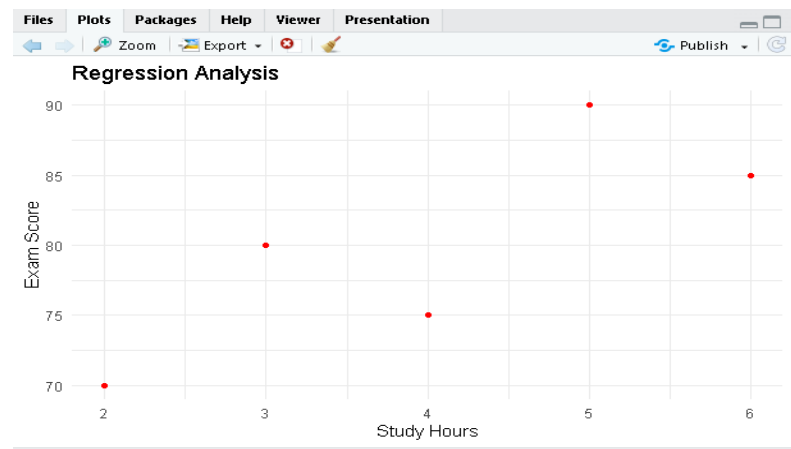
#### **# View the predicted marks**

```
print(predicted_marks)
```

## # Plot the data and the regression line

```
ggplot(Data_Frame, aes(x = study_hours, y = examscore)) +  
  geom_point(color = "Red") +  
  geom_line(data = new_data, aes(x = study_hours, y = predicted_marks), color = "red") +  
  labs(title = "Regression Analysis",  
        x = "Study Hours",  
        y = "Exam Score") +  
  theme_minimal()
```

## Output:





## 5. Linear Regression and Model Evaluation with Salary Data

### Dataset:

S.No.	Years of Experienced	Salary
1	1.1	39343
2	1.3	46205
3	1.5	37731
4	2.0	43525
5	2.2	39891
6	2.9	56642
7	3.0	60150
8	3.2	54445
9	3.2	64445
10	3.7	57189

### **Aim:**

To apply linear regression analysis to model salary data and evaluate the effectiveness of the model in predicting salaries based on relevant features.

### **Algorithm:**

Step 1: Install and Load Required Packages

- Install and load necessary R packages for data manipulation, visualization, and machine learning.

Step 2: Load Data

- Load your dataset "income.xlsx" using the "read\_excel" function from the "readxl" package.

Step 3: Create a Scatter Plot with Linear Regression Line

- Create a scatter plot with a linear regression line to visualize the relationship between "YearsExperienced" and "Salary."

Step 4: Data Splitting

- Split the data into a training set (70%) and a testing set (30%) based on the "Salary" column.

Step 5: Build a Linear Regression Model

- Fit a linear regression model to predict "Salary" based on "YearsExperienced" using the training data.

Step 6: Make Predictions and Evaluate

- Predict salaries on the test data and evaluate the model's performance.

Step 7: Additional Model Evaluation

**- Calculate Mean Squared Error (MSE) and R-squared for further evaluation.**

**Code:**

```
install.packages("readxl")
install.packages("caTools")
install.packages("dplyr")
install.packages("ggplot2")

library(readxl)
library(caTools)
library(readr)
library(dplyr)
library(ggplot2)

dataset1 <- read_excel("Directory of income.xlsx ")
dataset1

View(dataset1)

# Create the scatter plot with linear regression line
ggplot(dataset1, aes(x = YearsExperienced, y = Salary)) +
  geom_point() +
  geom_smooth(method = "lm", color = "blue") +
  labs(x = "Years Experienced", y = "Salary", title = "Scatter Plot of Years Experienced vs.
Salary with Linear Regression Line")

# Assuming you want to split the data into 70% for training and 30% for testing
split_ratio <- 0.7

# Perform the data split based on the "Salary" column
split <- caTools::sample.split(dataset1$Salary, SplitRatio = split_ratio)

# Extract the training and testing datasets based on the split
training_data <- subset(dataset1, split == TRUE)
testing_data <- subset(dataset1, split == FALSE)
```

### **# Fit the linear regression model on the training data**

```
lm_model <- lm(Salary ~ YearsExperienced, data = training_data)
```

### **# Predict the salaries on the test data**

```
predicted_salaries <- predict(lm_model, newdata = testing_data)
```

### **# Convert continuous predictions to discrete classes (High/Low) based on a threshold**

```
threshold <- 40000
```

```
predicted_classes <- ifelse(predicted_salaries > threshold, "High", "Low")
```

### **# Convert actual salaries to discrete classes (High/Low) based on the same threshold**

```
actual_classes <- ifelse(testing_data$Salary > threshold, "High", "Low")
```

### **# Calculate accuracy**

```
accuracy <- mean(predicted_classes == actual_classes)
```

### **# Calculate precision**

```
precision <- sum(predicted_classes == "High" & actual_classes == "High") /  
sum(predicted_classes == "High")
```

### **# Calculate recall**

```
recall <- sum(predicted_classes == "High" & actual_classes == "High") / sum(actual_classes  
== "High")
```

### **# Calculate F1 score**

```
f1_score <- 2 * precision * recall / (precision + recall)
```

### **# Display the results**

```
cat("Accuracy:", accuracy, "\n")
```

```
cat("Precision:", precision, "\n")
```

```
cat("Recall:", recall, "\n")
```

```
cat("F1 Score:", f1_score, "\n")
```

### **# Calculate the Mean Squared Error (MSE)**

```
mse <- mean((testing_data$Salary - predicted_salaries)^2)
```

### **# Calculate the R-squared**

```
rsquared <- summary(lm_model)$r.squared
```


## # Display the results

```
print(predicted_salaries)
```

```
cat("Mean Squared Error (MSE):", mse, "\n")
```

```
cat("R-squared:", rsquared, "\n")
```

## Output:



## 6. Alcohol Consumption Analysis

**Dataset:**

S. No.	User_ID	First_Name	Last_Name	Age	Gender	Location	Drinks_per_Week	Incidents_per_Month	Program_Status
1	1001	John	Doe	28	Male	City A	12	2	Enrolled
2	1002	Jane	Smith	35	Female	City B	5	1	Enrolled
3	1003	Alex	Lee	22	Male	City A	20	3	Not Enrolled
4	1004	Emily	Chen	42	Female	City C	8	0	Enrolled
5	1005	Mark	Brown	31	Male	City B	15	2	Not Enrolled

**Aim:**

To analyze patterns of alcohol consumption and its determinants, contributing to a better understanding of drinking behaviors and potential factors influencing them.

**Procedure:**

### Step 1: Load Necessary Libraries

- Begin by loading the required library for data visualization:

### Step 2: Create Sample Alcohol User Data

- Create a sample data frame named "alcohol\_data" containing information about alcohol users. Replace this with your actual data when working with your dataset:

### Step 3: Generate Summary Statistics

- Calculate summary statistics for key variables in the "alcohol\_data" data frame. For example:

### Step 4: Create a Scatter Plot

- Visualize the relationship between variables by creating a scatter plot using ggplot2. This step includes:

- Setting up data and aesthetics.
- Adding points to the plot.
- Labelling the title and axes.

**Code:**

**# Load necessary libraries**

library(ggplot2) # For data visualization

**# Sample alcohol user data (replace this with your actual data)**

```

alcohol_data <- data.frame(
  User_ID = c(1001, 1002, 1003, 1004, 1005),
  First_Name = c("John", "Jane", "Alex", "Emily", "Mark"),
  Last_Name = c("Doe", "Smith", "Lee", "Chen", "Brown"),
  Age = c(28, 35, 22, 42, 31),
  Gender = c("Male", "Female", "Male", "Female", "Male"),
  Location = c("City A", "City B", "City A", "City C", "City B"),
  Drinks_per_Week = c(12, 5, 20, 8, 15),
  Incidents_per_Month = c(2, 1, 3, 0, 2),
  Program_Status = c("Enrolled", "Enrolled", "Not Enrolled", "Enrolled", "Not Enrolled")
)

```

#### **# Summary statistics**

```

summary(alcohol_data$Drinks_per_Week)
summary(alcohol_data$Incidents_per_Month)

```

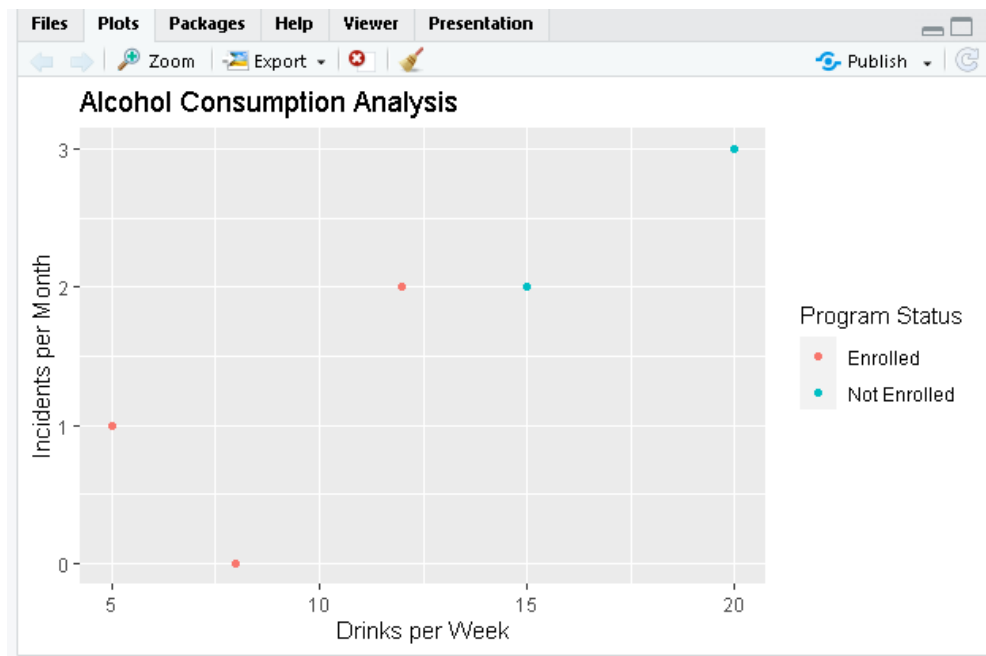
#### **# Create a scatter plot of Drinks per Week vs. Incidents per Month**

```

ggplot(alcohol_data, aes(x = Drinks_per_Week, y = Incidents_per_Month, color =
Program_Status)) +
  geom_point() +
  labs(title = "Alcohol Consumption Analysis",
    x = "Drinks per Week",
    y = "Incidents per Month",
    color = "Program Status")

```

#### **Output:**



---

---

---

## 7. Alcohol Consumption Analysis (Bar-Chart)

### Dataset:

S.No.	User_ID	First_Name	Last_Name	Age	Gender	Location	Drinks_per_Week	Incidents_per_Month	Program_Status
1	1001	John	Doe	28	Male	City A	12	2	Enrolled
2	1002	Jane	Smith	35	Female	City B	5	1	Enrolled
3	1003	Alex	Lee	22	Male	City A	20	3	Not Enrolled
4	1004	Emily	Chen	42	Female	City C	8	0	Enrolled
5	1005	Mark	Brown	31	Male	City B	15	2	Not Enrolled

### Aim:

To visually represent and analyze alcohol consumption data using bar charts, allowing for a clear and insightful depiction of drinking patterns and preferences.

### Algorithm:

Step 1: Load Necessary Libraries

- Start by loading the "ggplot2" library, which is used for data visualization.

Step 2: Create Sample Alcohol User Data

- Create a sample data frame named "alcohol\_data" with information about alcohol users. Replace this sample data with your actual data when working with real datasets.

Step 3: Calculate Average Drinks per Week by Program Status

- Use the `aggregate` function to calculate the average "Drinks\_per\_Week" for each "Program\_Status."

Step 4: Create a Bar Chart

- Create a bar chart to visualize the average "Drinks\_per\_Week" for each "Program\_Status."
- Use ggplot2 for this purpose.

### Code:

```
# Load necessary libraries
```

```
library(ggplot2) # For data visualization
```



**# Sample alcohol user data (replace this with your actual data)**

```
alcohol_data <- data.frame(  
  User_ID = c(1001, 1002, 1003, 1004, 1005),  
  First_Name = c("John", "Jane", "Alex", "Emily", "Mark"),  
  Last_Name = c("Doe", "Smith", "Lee", "Chen", "Brown"),  
  Age = c(28, 35, 22, 42, 31),  
  Gender = c("Male", "Female", "Male", "Female", "Male"),  
  Location = c("City A", "City B", "City A", "City C", "City B"),  
  Drinks_per_Week = c(12, 5, 20, 8, 15),  
  Incidents_per_Month = c(2, 1, 3, 0, 2),  
  Program_Status = c("Enrolled", "Enrolled", "Not Enrolled", "Enrolled", "Not Enrolled")  
)
```

**# Create a bar chart of average Drinks per Week by Program Status**

```
avg_drinks_by_status <- aggregate(Drinks_per_Week ~ Program_Status, data =  
alcohol_data, FUN = mean)  
  
ggplot(avg_drinks_by_status, aes(x = Program_Status, y = Drinks_per_Week, fill =  
Program_Status)) +  
  geom_bar(stat = "identity") +  
  labs(title = "Average Drinks per Week by Program Status",  
    x = "Program Status",  
    y = "Average Drinks per Week") +  
  theme_minimal() +  
  theme(legend.position = "none")
```



## Output:

